

INDEX

SR. NO.	TOPIC	DATE	SIGNATURE
1	a. Perform conditional formatting on a datasheet using various criteria b. Creative pivot table to analyse and summarise data c. Use VLOOKUP function to retrieve information from a different worksheet or a table d. Perform what if analysis using goal seek to determine input values for desired output		
2.	Data frames and basic data pre processing		
3.	Feature scaling and dummification		
4.	Hypothesis testing		
5.	ANOVA (Analysis of Variance)		
6.	Regression and its types		
7.	Logistic regression and decision tree		
8.	K means clustering		
9.	Principal component analysis PCA		
10.	Data visualisation and story telling		

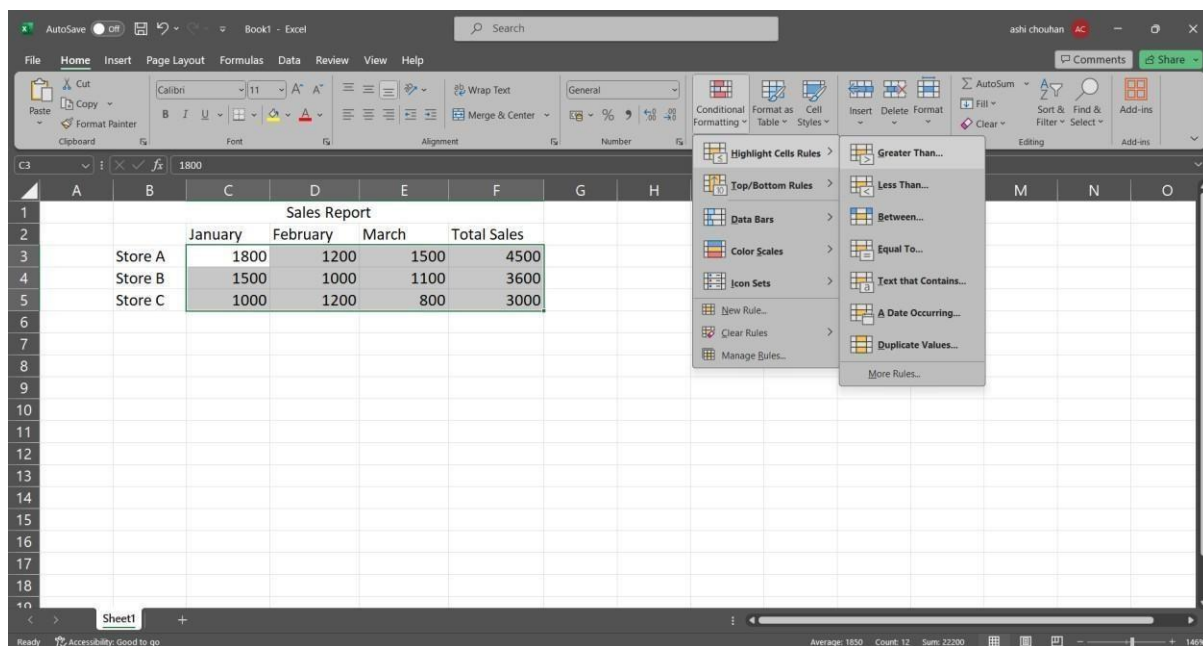
PRACTICAL NO. – 1

A. Perform conditional formatting on a dataset using various criteria.

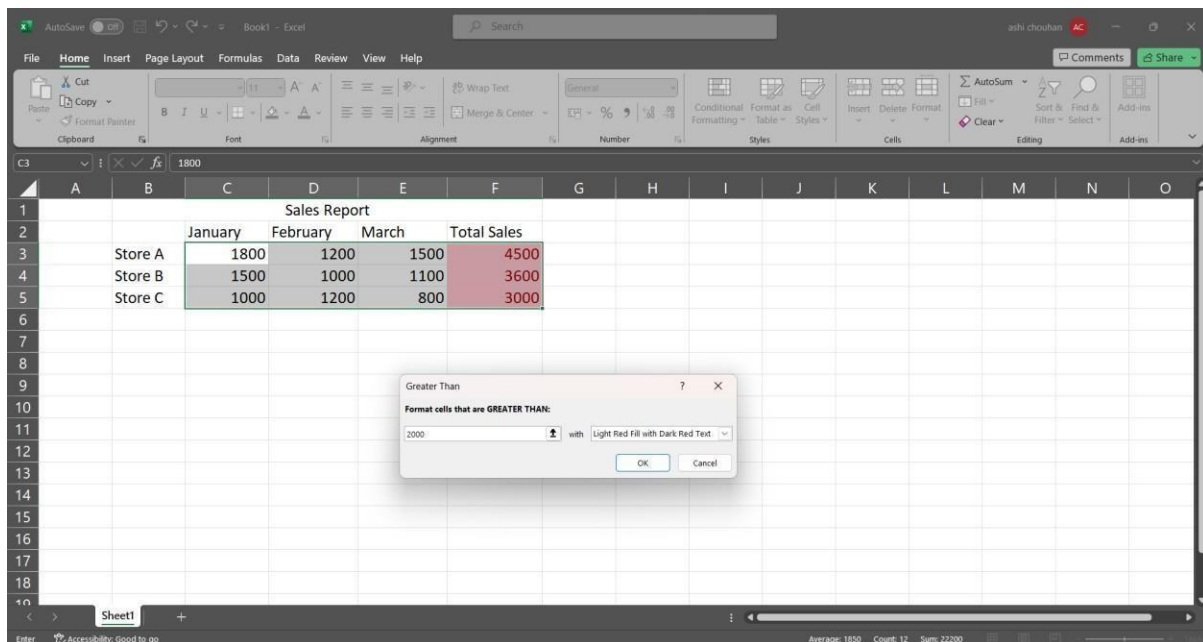
	A	B	C	D	E	F	G	H
1								
2								
3			January	February	March	Total Sales		
4		Store A	1800	1200	1500	4500		
5		Store B	1500	1000	1100	3600		
6		Store C	1000	1200	800	3000		
7								
8								
9								
10								
11								
12								

Steps

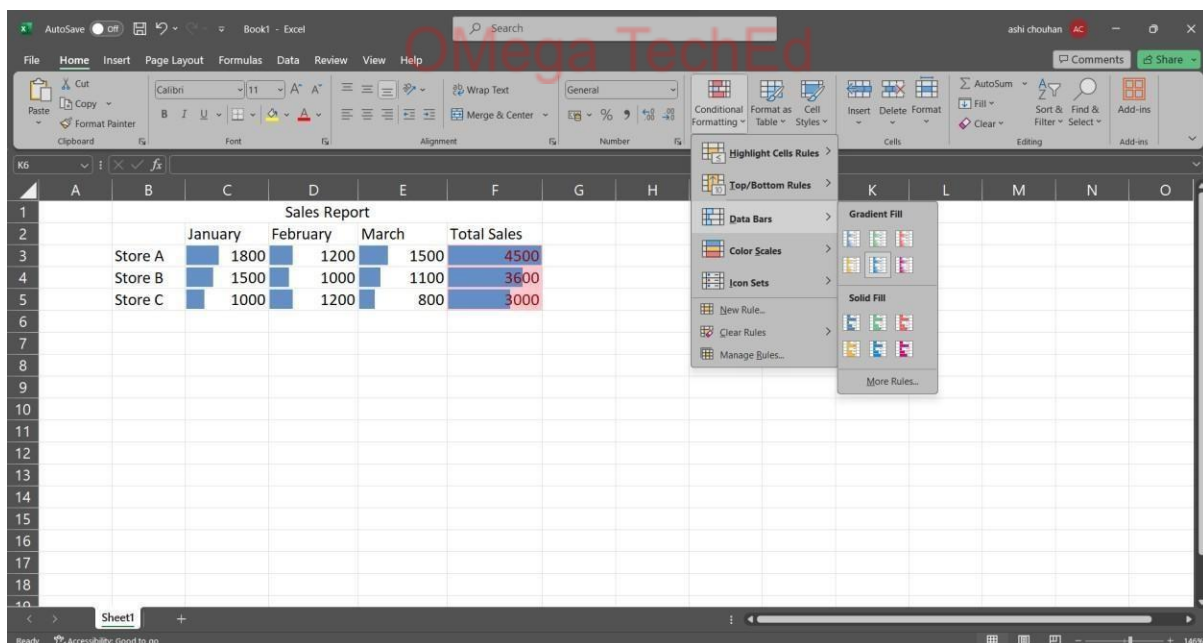
Step 1: Go to conditional formatting > Greater Than



Step 2: Enter the greater than filter value for example 2000.



Step 3: Go to Data Bars > Solid Fill in conditional formatting.

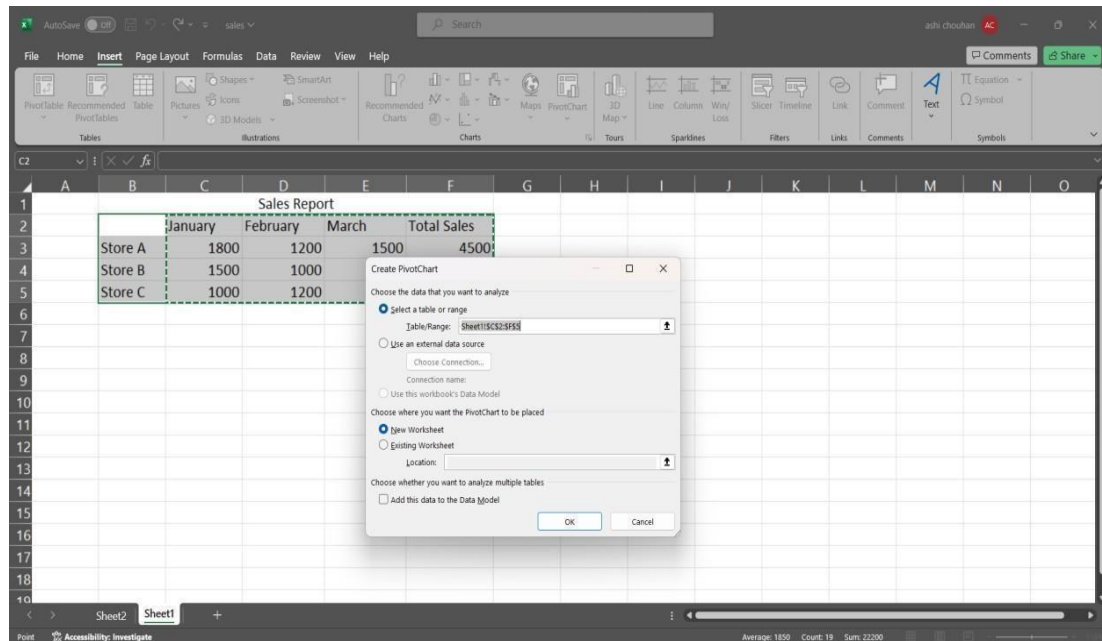


B. Create a pivot table to analyse and summarize data.

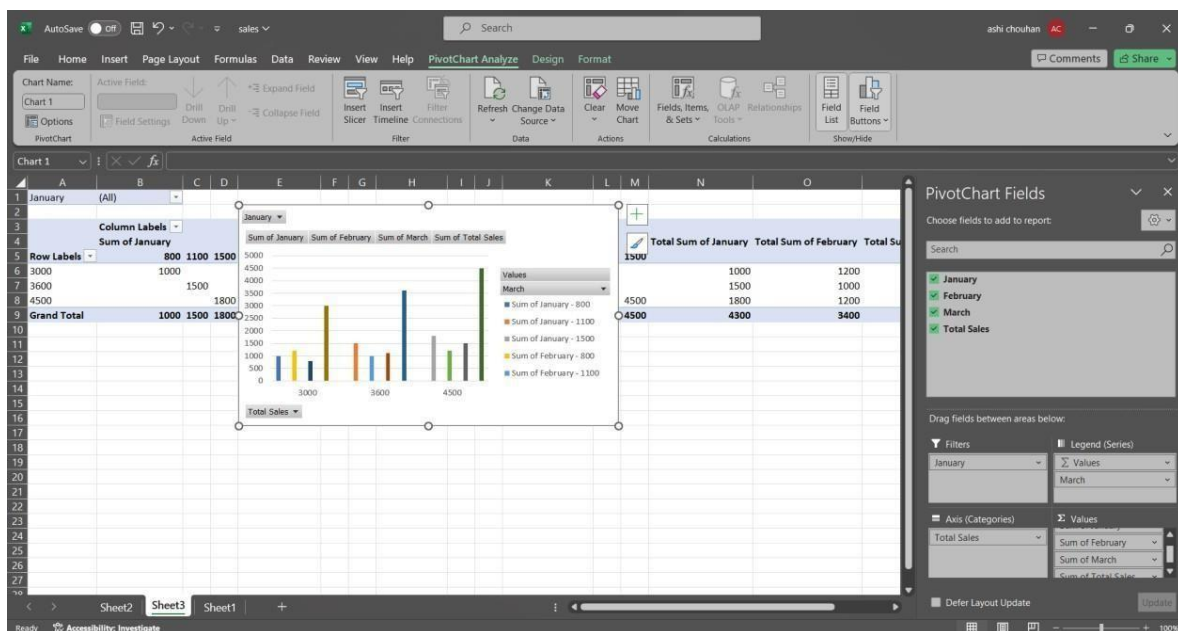
Steps

Step1: - Select the entire table and go to Insert Tab PivotChart > Pivotchart

Step2:- Select “New Worksheet” in the create pivot chart window.



Step 3: Select and drag attributes in the below boxes.

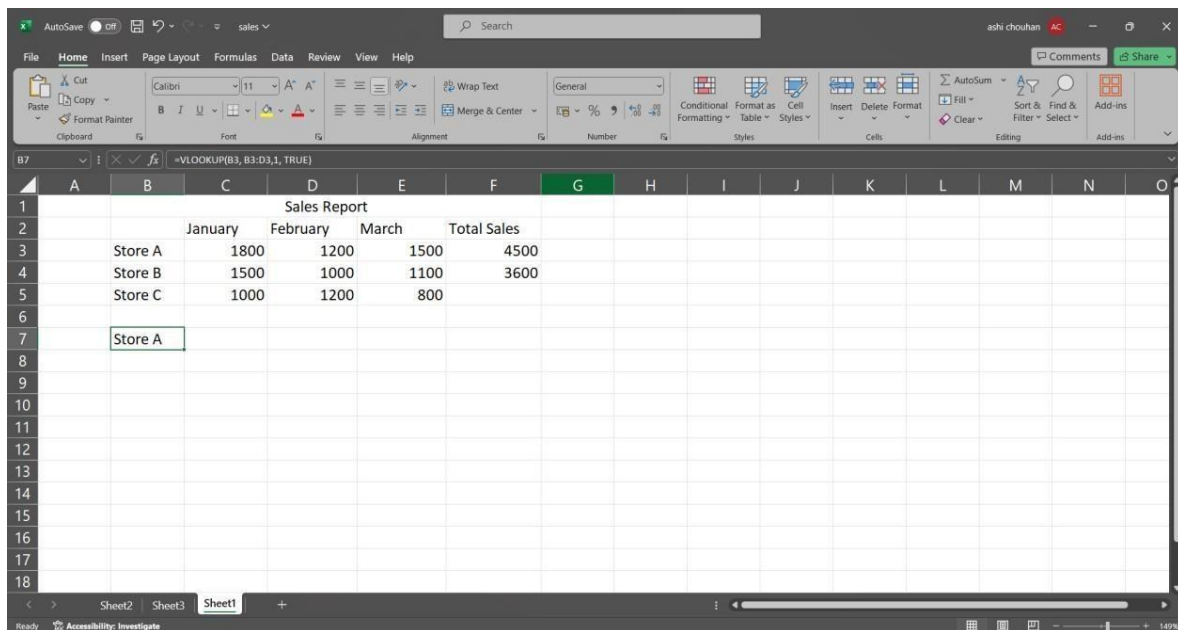


C. Use VLOOKUP function to retrieve information from a different worksheet or table.

Steps

Step 1: click on an empty cell and type the following command.

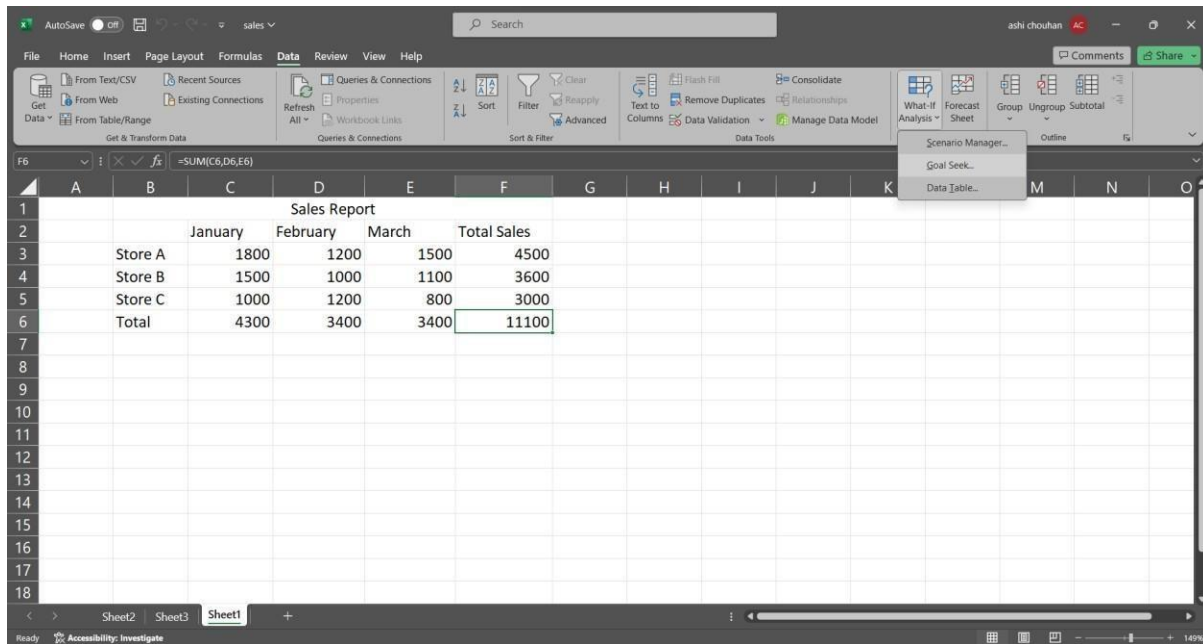
=VLOOKUP(B3, B3:D3,1, TRUE)



D. Perform what-if analysis using Goal Seek to determine input values for desired output.

Steps

Step 1: In the Data tab go to the what if analysis>Goal seek.



Step 2: Fill the information in the window accordingly and click ok.

The screenshot displays the Microsoft Excel interface with a 'Sales Report' table and a 'Goal Seek Status' dialog box.

Sales Report Table:

	January	February	March	Total Sales
Store A	1800	1200	1500	4500
Store B	1500	1000	1100	3600
Store C	1000	1200	109700	111900
Total	4300	3400	112300	120000

Goal Seek Status Dialog Box:

Goal Seeking with Cell F6 found a solution.

Target value: 120000
Current value: 120000

Buttons: Step, Pause, OK, Cancel

The screenshot shows the Microsoft Excel interface with the 'Data' tab selected. A table titled 'Sales Report' is displayed in the range A2:F6. The table has columns for months (January, February, March) and Total Sales. The formula bar shows the formula for cell F6: `=SUM(C6,D6,E6)`. The 'Goal Seek' dialog box is open, with the following settings:

- Set cell: F6
- To value: 12000
- By changing cell: E5

The dialog box has 'OK' and 'Cancel' buttons.

	January	February	March	Total Sales
Store A	1800	1200	1500	4500
Store B	1500	1000	1100	3600
Store C	1000	1200	800	3000
Total	4300	3400	3400	11100

The screenshot shows the same Excel interface, but the 'Goal Seek' dialog box is no longer present. The 'Sales Report' table has been updated with the following values:

	January	February	March	Total Sales
Store A	1800	1200	1500	4500
Store B	1500	1000	1100	3600
Store C	1000	1200	109700	111900
Total	4300	3400	112300	120000

PRACTICAL NO. 2

AIM:- Data Frames and Basic Data Pre-Processing

A. Read Data from CSV and JSON files into a data frame.

SOURCE CODE:-

```
#Read data from a csv file
import pandas as pd
df=pd.read_csv('Student_Marks.csv')
print("Our dataset")
print(df)
```

OUTPUT:-

```
===== RESTART: D:\Notes\sem-6\data science\prac2
Our dataset
   number_courses  time_study  Marks
0                3      4.508  19.202
1                4      0.096   7.734
2                4      3.133  13.811
3                6      7.909  53.018
4                8      7.811  55.299
..             ...         ...    ...
95                6      3.561  19.128
96                3      0.301   5.609
97                4      7.163  41.444
98                7      0.309  12.027
99                3      6.335  32.357

[100 rows x 3 columns]
>>>
```

SOURCE CODE:-

```
#Reading data from a JSON Import
pandas as pd
data=pd.read_json('dataset.json')
print(data)
```

OUTPUT:-

```
>>> ===== RESTART: D:/Notes/sem-6/data science/p:
      fruit      size  color
0   Apple   Large    Red
1  Banana  Medium  Yellow
2  Orange   Small  Orange
>>>
```

B. Perform basic data pre-processing tasks such as handling missing values and outliers code:

1. #Replacing NA values using fillna()

```
Import pandas as pd
df=pd.read_csv('titanic.csv')
Print(df)
df.head(10)
print("Dataset after filling NA values with 0:") df2 =
df.fillna(value = 0)
print(df2)
```

OUTPUT:-

OMega TechEd

```

===== RESTART: D:/Notes/sem-6/data science/prac2c.py =====
      PassengerId  Survived  Pclass    ...    Cabin Embarked
0              892         0.0    3.0    ...   NaN      Q
1              893         0.0    3.0    ...   NaN      S
2              894         0.0    2.0    ...   NaN      Q
3              895         0.0    3.0    ...   NaN      S
4              896         0.0    3.0    ...   NaN      S
..              ...         ...    ...    ...   ...      ...
413            1305         0.0    3.0    ...   NaN      S
414            1306         0.0    1.0    ...  C105      C
415            1307         0.0    3.0    ...   NaN      S
416            1308         0.0    3.0    ...   NaN      S
417            1309         0.0    3.0    ...   NaN      C

[418 rows x 11 columns]
Dataset after filling NA values with 0 :
      PassengerId  Survived  Pclass    ...    Cabin Embarked
0              892         0.0    3.0    ...     0      Q
1              893         0.0    3.0    ...     0      S
2              894         0.0    2.0    ...     0      Q
3              895         0.0    3.0    ...     0      S
4              896         0.0    3.0    ...     0      S
..              ...         ...    ...    ...   ...      ...
413            1305         0.0    3.0    ...     0      S
414            1306         0.0    1.0    ...  C105      C
415            1307         0.0    3.0    ...     0      S
416            1308         0.0    3.0    ...     0      S
417            1309         0.0    3.0    ...     0      C

[418 rows x 11 columns]
.>>

```

2. #Dropping NA values using dropna()

Import pandas as pd

```
df=pd.read_csv('titanic.csv')
```

```
print(df)
```

```
df.head(10)
```

```
print("Dataset after dropping NA values:") df.dropna(inplace
= True)
```

```
print(df)
```

OUTPUT:-

```

===== RESTART: D:/Notes/sem-6/data science/prac2c.py =====
   PassengerId  Survived  Pclass    Cabin Embarked
0            892         3.0    NaN      Q
1            893         3.0    NaN      S
2            894         2.0    NaN      Q
3            895         3.0    NaN      S
4            896         NaN    NaN      S
..         ...         ...    ...    ...
413          1305         3.0    NaN      S
414          1306         1.0    C105      C
415          1307         3.0    NaN      S
416          1308         3.0    NaN      S
417          1309         3.0    NaN      C

[418 rows x 11 columns]
Dataset after dropping NA values:
   PassengerId  Survived  Pclass    Cabin Embarked
12            904         1.0    NaN      B45      S
14            906         1.0    NaN      E31      S
24            916         1.0    B57 B59 B63 B66      C
26            918         1.0    NaN      B36      C
28            920         1.0    NaN      A21      S
..         ...         ...    ...    ...
404          1296         1.0    NaN      D40      C
405          1297         2.0    NaN      D38      C
407          1299         1.0    NaN      C80      C
411          1303         1.0    NaN      C78      Q
414          1306         1.0    C105      C

[87 rows x 11 columns]
>>>

```

C. Manipulate and transform data using functions like

filtering, sorting, and grouping Code:

SOURCE CODE:-

```
import pandas as pd #
```

Load iris dataset

```
iris = pd.read_csv('Iris.csv')
```

Filtering data based on a

```
condition setosa = iris[iris['Species'] ==
```

```
'setosa'] print("Setosa samples:")
```

```
print(setosa.head())
```

Sorting data

```
sorted_iris =
```

```
iris.sort_values(by='SepalLengthCm',ascending=
```

```
g=False)

print("\nSorted iris dataset:")

print(sorted_iris.head())

# Grouping data grouped_species =
iris.groupby('Species').mean() print("\nMean
measurements for each species:")
print(grouped_species)
```

OUTPUT:-

```
===== RESTART: D:/Notes/sem-6/data science/prac2b.py =====
Setosa samples:
Empty DataFrame
Columns: [Id, SepalLengthCm, SepalWidthCm, PetalLengthCm, PetalWidthCm, Species]
Index: []

Sorted iris dataset:
   Id  SepalLengthCm  ...  PetalWidthCm  Species
131 132           7.9  ...           2.0  Iris-virginica
135 136           7.7  ...           2.3  Iris-virginica
122 123           7.7  ...           2.0  Iris-virginica
117 118           7.7  ...           2.2  Iris-virginica
118 119           7.7  ...           2.3  Iris-virginica

[5 rows x 6 columns]

Mean measurements for each species:
      Id  SepalLengthCm  ...  PetalLengthCm  PetalWidthCm
Species
Iris-setosa      25.5      5.006  ...      1.464      0.244
Iris-versicolor  75.5      5.936  ...      4.260      1.326
Iris-virginica  125.5      6.588  ...      5.552      2.026

[3 rows x 5 columns]
>>
```

PRACTICAL NO. 3

AIM:- Feature Scaling and Dummification

- A. Apply feature-scaling techniques like standardization and normalization to numerical features.

SOURCE CODE:-

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler,
StandardScaler #
Load dataset
df = pd.read_csv('wine.csv', header=None, usecols=[0, 1,
2], skiprows=1)
df.columns = ['classlabel', 'Alcohol', 'Malic Acid'] #
Display original DataFrame
print("Original DataFrame:")
print(df)
# Apply Min-Max Scaling min_max_scaler
= MinMaxScaler() df[['Alcohol', 'Malic
Acid']] =
min_max_scaler.fit_transform(df[['Alcohol', 'Malic Acid']])
print("\nDataFrame after Min-Max Scaling:") print(df)
# Apply Standard Scaling standard_scaler =
StandardScaler() df[['Alcohol', 'Malic Acid']]
=
standard_scaler.fit_transform(df[['Alcohol', 'Malic Acid']])
print("\nDataFrame after Standard Scaling:") print(df)
```

OUTPUT:-

```

= RESTART: D:/Notes/sem-6/data science/prac3b.py
Original DataFrame:
  classlabel  Alcohol  Malic Acid
0          1    14.23      1.71
1          1    13.20      1.78
2          1    13.16      2.36
3          1    14.37      1.95
4          1    13.24      2.59
..         ...      ...      ...
173        3    13.71      5.65
174        3    13.40      3.91
175        3    13.27      4.28
176        3    13.17      2.59
177        3    14.13      4.10

[178 rows x 3 columns]

Dataframe after MinMax Scaling
  classlabel  Alcohol  Malic Acid
0          1  0.842105  0.191700
1          1  0.571053  0.205534
2          1  0.560526  0.320158
3          1  0.878947  0.239130
4          1  0.581579  0.365613
..         ...      ...      ...
173        3  0.705263  0.970356
174        3  0.623684  0.626482
175        3  0.589474  0.699605
176        3  0.563158  0.365613
177        3  0.815789  0.664032

[178 rows x 3 columns]

Dataframe after Standard Scaling
  classlabel  Alcohol  Malic Acid
173        3  0.303019  0.303019
177        3  0.815789  0.664032

[178 rows x 3 columns]

Dataframe after Standard Scaling
  classlabel  Alcohol  Malic Acid
0          1  1.518613 -0.562250
1          1  0.246290 -0.499413
2          1  0.196879  0.021231
3          1  1.691550 -0.346811
4          1  0.295700  0.227694
..         ...      ...      ...
173        3  0.876275  2.974543
174        3  0.493343  1.412609
175        3  0.332758  1.744744
176        3  0.209232  0.227694
177        3  1.395086  1.583165

[178 rows x 3 columns]
>>>

```

- B. Perform feature Dummification to convert categorical variables into numerical representations.

SOURCE CODE:-

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder # Load
the dataset
iris = pd.read_csv("Iris.csv") print(iris)
# Apply Label Encoding le =
LabelEncoder()
iris['code'] = le.fit_transform(iris['Species'])
print(iris)
```

OUTPUT:-

```
===== RESTART: D:/Notes/sem-6/data science/prac3a.py =====
   Id  SepalLengthCm  ...  PetalWidthCm  Species
0    1           5.1  ...           0.2  Iris-setosa
1    2           4.9  ...           0.2  Iris-setosa
2    3           4.7  ...           0.2  Iris-setosa
3    4           4.6  ...           0.2  Iris-setosa
4    5           5.0  ...           0.2  Iris-setosa
..  ...           ...  ...           ...  ...
145 146           6.7  ...           2.3  Iris-virginica
146 147           6.3  ...           1.9  Iris-virginica
147 148           6.5  ...           2.0  Iris-virginica
148 149           6.2  ...           2.3  Iris-virginica
149 150           5.9  ...           1.8  Iris-virginica

[150 rows x 6 columns]
   Id  SepalLengthCm  SepalWidthCm  ...  PetalWidthCm  Species  code
0    1           5.1           3.5  ...           0.2  Iris-setosa    0
1    2           4.9           3.0  ...           0.2  Iris-setosa    0
2    3           4.7           3.2  ...           0.2  Iris-setosa    0
3    4           4.6           3.1  ...           0.2  Iris-setosa    0
4    5           5.0           3.6  ...           0.2  Iris-setosa    0
..  ...           ...           ...  ...           ...  ...  ...
145 146           6.7           3.0  ...           2.3  Iris-virginica    2
146 147           6.3           2.5  ...           1.9  Iris-virginica    2
147 148           6.5           3.0  ...           2.0  Iris-virginica    2
148 149           6.2           3.4  ...           2.3  Iris-virginica    2
149 150           5.9           3.0  ...           1.8  Iris-virginica    2

[150 rows x 7 columns]
>>>
```


PRACTICAL NO. – 4**AIM:- Hypothesis Testing****SOURCE CODE:-**

```
import numpy as np from
scipy import stats
import matplotlib.pyplot as plt

# Generate two samples for demonstration purposes np.random.seed(42)
sample1 = np.random.normal(loc=10, scale=2, size=30) sample2 =
np.random.normal(loc=12, scale=2, size=30)

# Perform a two-sample t-test
t_statistic, p_value = stats.ttest_ind(sample1, sample2)

# Set the significance level alpha
= 0.05

print("Results of Two-Sample t-test:")
print(f'T-statistic: {t_statistic}') print(f'P-
value: {p_value}')
print(f'Degrees of Freedom: {len(sample1) + len(sample2) - 2}')
```

```
# Plot the distributions
```

```
plt.figure(figsize=(10, 6))
```

```
plt.hist(sample1, alpha=0.5, label='Sample 1', color='blue')
```

```
plt.hist(sample2, alpha=0.5, label='Sample 2', color='orange')
```

```
plt.axvline(np.mean(sample1), color='blue', linestyle='dashed',  
linewidth=2)
```

```
plt.axvline(np.mean(sample2), color='orange', linestyle='dashed', linewidth=2)
```

```
plt.title('Distributions of Sample 1 and Sample 2') plt.xlabel('Values')
```

```
plt.ylabel('Frequency')
```

```
plt.legend()
```

```
# Highlight the critical region if null hypothesis is rejected if
```

```
p_value < alpha:
```

```
    critical_region = np.linspace(min(sample1.min(), sample2.min()),  
max(sample1.max(), sample2.max()), 1000)
```

```
    plt.fill_between(critical_region, 0, 5, color='red', alpha=0.3,  
label='Critical Region')
```

```
    plt.text(11, 5, f'T-statistic: {t_statistic:.2f}', ha='center', va='center', color='black',  
backgroundcolor='white')
```

```
# Show the plot
```

```
plt.show()
```

```
# Draw Conclusions if
```

```
p_value < alpha:
```

```
    if np.mean(sample1) > np.mean(sample2):
```

```
        print("Conclusion: There is significant evidence to reject the null hypothesis.")
```

```
        print("Interpretation: The mean of Sample 1 is significantly higher than  
that of Sample 2.")
```

```
    else:
```

```
        print("Conclusion: There is significant evidence to reject the null hypothesis.")
```

```
        print("Interpretation: The mean of Sample 2 is significantly higher than  
that of Sample 1.")
```


```
else:
```

```
    print("Conclusion: Fail to reject the null hypothesis.") print("Interpretation:
```

```
    There is not enough evidence to claim a  
significant difference between the means.")
```

OUTPUT:-

```
----- RESTART: In / all hou  
Results of Two-Sample t-test:  
T-statistic: -4.512913234547555  
P-value: 3.176506547470154e-05  
Degrees of Freedom: 58
```



SOURCE CODE:-

```
import numpy as
np import pandas
as pd
import matplotlib.pyplot as plt
import seaborn as sb
import warnings
from scipy import stats

# Suppress warnings
warnings.filterwarnings('ignore')

# Load dataset
df = sb.load_dataset('mpg')
```

```
print(df)
```

```
# Describe horsepower and model year columns
```

```
print(df['horsepower'].describe())
```

```
print(df['model_year'].describe())
```

```
# Categorize horsepower into bins
```

```
bins = [0, 75, 150, 240]
```

```
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm',
```

```
'h']) c = df['horsepower_new']
```

```
print(c)
```

```
# Categorize model year into bins
```

```
ybins = [69, 72, 74, 84]
```

```
labels = ['t1', 't2', 't3']
```

```
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins, labels=labels)
```

```
newyear = df['modelyear_new']
```

```
print(newyear)
```

```
# Create a contingency table
```

```
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new'])
```

```
print(df_chi)
```

```
# Perform chi-square test
```

```
chi2_stat, p_value, dof, expected =
```

```
stats.chi2_contingency(df_chi) print(f'Chi-Square Statistic:  
{chi2_stat}')  
  
print(f'P-value: {p_value}')  
  
print(f'Degrees of Freedom: {dof}')  
  
print(f'Expected Frequencies:\n{expected}')
```

OUTPUT:-

```

----- RESIDENT: E:/all notes/DS/plac_4.1.py -----
      mpg  cylinders  ...  origin  name
0      18.0         8  ...    usa  chevrolet chevelle malibu
1      15.0         8  ...    usa  buick skylark 320
2      18.0         8  ...    usa  plymouth satellite
3      16.0         8  ...    usa  amc rebel sst
4      17.0         8  ...    usa  ford torino
..      ...         ...  ...    ...  ...
393    27.0         4  ...    usa  ford mustang gl
394    44.0         4  ...  europe  vw pickup
395    32.0         4  ...    usa  dodge rampage
396    28.0         4  ...    usa  ford ranger
397    31.0         4  ...    usa  chevy s-10

[398 rows x 9 columns]
count      392.000000
mean       104.469388
std        38.491160
min         46.000000
25%         75.000000
50%         93.500000
75%        126.000000
max        230.000000

```

```

Name: horsepower, dtype: float64
count      398.000000
mean         76.010050
std          3.697627
min          70.000000
25%          73.000000
50%          76.000000
75%          79.000000
max          82.000000
Name: model_year, dtype: float64
0          m
1          h
2          m
3          m
4          m
..
393        m
394         l
395        m
396        m
397        m

```

OMega TechEd

```
Name: horsepower_new, Length: 398, dtype: category
Categories (3, object): ['l' < 'm' < 'h']
0      t1
1      t1
2      t1
3      t1
4      t1
..
393    t3
394    t3
395    t3
396    t3
397    t3
Name: modelyear_new, Length: 398, dtype: category
Categories (3, object): ['t1' < 't2' < 't3']
modelyear_new  t1  t2  t3
horsepower_new
l              9  14  76
m             49  41 158
h             26  11   8
(54.95485392447537, 3.320518009555984e-11, 4, array([[ 21.21428571,  16.66836735,  61.11734694]
,
          [ 53.14285714,  41.75510204, 153.10204082],
          [  9.64285714,   7.57653061,  27.78061224]])
```


PRACTICAL NO. – 5**AIM:- ANOVA(Analysis of Variance)****SOURCE CODE:-**

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import warnings
from scipy import stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd

# Suppress warnings
warnings.filterwarnings('ignore')

# Load dataset
df = sb.load_dataset('mpg')
print(df)

# Describe horsepower and model year columns
print(df['horsepower'].describe())
print(df['model_year'].describe())

# Categorize horsepower into bins
bins = [0, 75, 150, 240]
```

```
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm', 'h'])
c = df['horsepower_new'] print(c)
```

```
# Categorize model year into bins ybins =
[69, 72, 74, 84]
labels = ['t1', 't2', 't3']
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins,
labels=labels)
newyear = df['modelyear_new']
print(newyear)
```

```
# Create a contingency table
```

```
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new']) print(df_chi)
```

```
# Perform chi-square test
```

```
chi2_stat, p_value, dof, expected = stats.chi2_contingency(df_chi)
```

```
print(f'Chi-Square Statistic: {chi2_stat}')
```

```
print(f'P-value: {p_value}') print(f'Degrees of
```

```
Freedom: {dof}')
```

```
print(f'Expected Frequencies:\n{expected}')
```

```
# Define groups for ANOVA
group1 = [23, 25, 29, 34, 30]
group2 = [19, 20, 22, 24, 25]
group3 = [15, 18, 20, 21, 17]
group4 = [28, 24, 26, 30, 29]

# Combine data into a DataFrame
data = pd.DataFrame({'value': group1 + group2 + group3 + group4, 'group':
                    ['Group1'] * len(group1) + ['Group2'] *
len(group2) +
                    ['Group3'] * len(group3) + ['Group4'] *
len(group4)})

# Perform one-way ANOVA
f_statistics, p_value_anova = stats.f_oneway(group1, group2, group3, group4)
print("\nOne-way ANOVA:")
print("F-statistics:", f_statistics)
print("P-value:", p_value_anova)

# Perform Tukey-Kramer post-hoc test
tukey_results = pairwise_tukeyhsd(data['value'], data['group'])
print("\nTukey-Kramer post-hoc test:")
print(tukey_results)
```

OUTPUT:-

```

=====
one-way ANOVA:
F-statistics: 12.139872842870115
p-value 0.00021465200901629603

Tukey-Kramer post-hoc test:
Multiple Comparison of Means - Tukey HSD, FWER=0.05
=====
group1 group2 meandiff p-adj  lower  upper  reject
-----
Group1 Group2      -6.2  0.024 -11.6809 -0.7191   True
Group1 Group3     -10.0 0.0004 -15.4809 -4.5191   True
Group1 Group4      -0.8 0.9747  -6.2809  4.6809  False
Group2 Group3      -3.8 0.2348  -9.2809  1.6809  False
Group2 Group4       5.4 0.0542  -0.0809 10.8809  False
Group3 Group4       9.2  0.001   3.7191 14.6809   True
=====
|

```

PRACTICAL NO. – 6**AIM:- Regression and its Types SOURCE****CODE:-**

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import warnings
from scipy import stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Suppress warnings
warnings.filterwarnings('ignore')

# Load dataset
df = sb.load_dataset('mpg')
print(df)

# Describe horsepower and model year columns
```

```
print(df['horsepower'].describe()) print(df['model_year'].describe())

# Categorize horsepower into bins bins =
[0, 75, 150, 240]
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm', 'h'])
c = df['horsepower_new'] print(c)

# Categorize model year into bins ybins =
[69, 72, 74, 84]
labels = ['t1', 't2', 't3']
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins,
labels=labels)
newyear = df['modelyear_new']
print(newyear)

# Create a contingency table
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new']) print(df_chi)

# Perform chi-square test
chi2_stat, p_value, dof, expected = stats.chi2_contingency(df_chi)
```

```
print(f'Chi-Square Statistic: {chi2_stat}')
print(f'P-value: {p_value}') print(f'Degrees of
Freedom: {dof}')
print(f'Expected Frequencies:\n{expected}')

# Define groups for ANOVA
group1 = [23, 25, 29, 34, 30]
group2 = [19, 20, 22, 24, 25]
group3 = [15, 18, 20, 21, 17]
group4 = [28, 24, 26, 30, 29]

# Combine data into a DataFrame
data = pd.DataFrame({'value': group1 + group2 + group3 + group4, 'group':
                        ['Group1'] * len(group1) + ['Group2'] *
len(group2) +
                        ['Group3'] * len(group3) + ['Group4'] *
len(group4)})

# Perform one-way ANOVA
f_statistics, p_value_anova = stats.f_oneway(group1, group2, group3, group4)
print("\nOne-way ANOVA:")
print("F-statistics:", f_statistics)
print("P-value:", p_value_anova)
```

```
# Perform Tukey-Kramer post-hoc test
tukey_results = pairwise_tukeyhsd(data['value'], data['group'])
print("\nTukey-Kramer post-hoc test:")
print(tukey_results)

# Load California housing dataset
housing = fetch_california_housing()
housing_df = pd.DataFrame(housing.data,
                           columns=housing.feature_names)
print(housing_df)

# Add target variable housing_df['PRICE'] =
housing.target

# Select feature and target
X = housing_df[['AveRooms']] y =
housing_df['PRICE']

# Split dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train linear regression model model =
LinearRegression()
```



```
model.fit(X_train, y_train)
```

```
# Evaluate model
```

```
mse = mean_squared_error(y_test, model.predict(X_test)) r2 =
```

```
r2_score(y_test, model.predict(X_test))
```

```
print("Mean Squared Error:", mse)
```

```
print("R-squared:", r2) print("Intercept:",
```

```
model.intercept_) print("Coefficient:",
```

```
model.coef_) OUTPUT:-
```

```
===
      MedInc HouseAge AveRooms ... AveOccup Latitude Longitude
0      8.3252    41.0  6.984127 ...   2.555556    37.88   -122.23
1      8.3014    21.0  6.238137 ...   2.109842    37.86   -122.22
2      7.2574    52.0  8.288136 ...   2.802260    37.85   -122.24
3      5.6431    52.0  5.817352 ...   2.547945    37.85   -122.25
4      3.8462    52.0  6.281853 ...   2.181467    37.85   -122.25
...
20635  1.5603    25.0  5.045455 ...   2.560606    39.48   -121.09
20636  2.5568    18.0  6.114035 ...   3.122807    39.49   -121.21
20637  1.7000    17.0  5.205543 ...   2.325635    39.43   -121.22
20638  1.8672    18.0  5.329513 ...   2.123209    39.43   -121.32
20639  2.3886    16.0  5.254717 ...   2.616981    39.37   -121.24

[20640 rows x 8 columns]
Mean Squared Error: 1.2923314440807299
R-squared: 0.013795337532284901
Intercept: 1.654762268596842
Coefficient: [0.07675559]
>
```

SOURCE CODE:-

```
import numpy as np import
```

```
pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sb
import warnings
from scipy import stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Suppress warnings
warnings.filterwarnings('ignore')

# Load dataset
df = sb.load_dataset('mpg')
print(df)

# Describe horsepower and model year columns
print(df['horsepower'].describe())
print(df['model_year'].describe())

# Categorize horsepower into bins
bins = [0, 75, 150, 240]
```

```
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm', 'h'])  
c = df['horsepower_new'] print(c)
```

```
# Categorize model year into bins ybins =  
[69, 72, 74, 84]  
labels = ['t1', 't2', 't3']  
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins, labels=labels)  
newyear = df['modelyear_new']  
print(newyear)
```

```
# Create a contingency table  
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new'])  
print(df_chi)
```

```
# Perform chi-square test  
chi2_stat, p_value, dof, expected = stats.chi2_contingency(df_chi)  
print(f'Chi-Square Statistic: {chi2_stat}') print(f'P-  
value: {p_value}')
```

```
print(f'Degrees of Freedom: {dof}') print(f'Expected  
Frequencies:\n{expected}')
```

```
# Define groups for ANOVA group1 = [23,  
25, 29, 34, 30]
```

```
group2 = [19, 20, 22, 24, 25]
```

```
group3 = [15, 18, 20, 21, 17]
```

```
group4 = [28, 24, 26, 30, 29]
```

```
# Combine data into a DataFrame
```

```
data = pd.DataFrame({'value': group1 + group2 + group3 + group4,  
                     'group': ['Group1'] * len(group1) + ['Group2'] * len(group2) +  
                               ['Group3'] * len(group3) + ['Group4'] *  
len(group4)}))
```

```
# Perform one-way ANOVA
```

```
f_statistics, p_value_anova = stats.f_oneway(group1, group2, group3, group4)  
print("\nOne-way ANOVA:") print("F-  
statistics:", f_statistics) print("P-value:",  
p_value_anova)
```

```
# Perform Tukey-Kramer post-hoc test

tukey_results = pairwise_tukeyhsd(data['value'], data['group']) print("\nTukey-
Kramer post-hoc test:")
print(tukey_results)


# Load California housing dataset housing =
fetch_california_housing()
housing_df = pd.DataFrame(housing.data, columns=housing.feature_names)
print(housing_df)


# Add target variable housing_df['PRICE'] =
housing.target


# Multiple Linear Regression

X = housing_df.drop('PRICE', axis=1) y =
housing_df['PRICE']


# Split dataset into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# Train linear regression model model =  
LinearRegression() model.fit(X_train,  
y_train)  
  
# Make predictions  
y_pred = model.predict(X_test)  
  
# Evaluate model  
mse = mean_squared_error(y_test, y_pred) r2 =  
r2_score(y_test, y_pred)  
print("Mean Squared Error:", mse) print("R-  
squared:", r2) print("Intercept:",  
model.intercept_) print("Coefficient:",  
model.coef_)
```

OUTPUT:-

```

----- RESIDUALS: E:\all notes\DS\prac_0_single.py -----
      MedInc  HouseAge  AveRooms  ...  AveOccup  Latitude  Longitude
0      8.3252     41.0    6.984127  ...    2.555556     37.88     -122.23
1      8.3014     21.0    6.238137  ...    2.109842     37.86     -122.22
2      7.2574     52.0    8.288136  ...    2.802260     37.85     -122.24
3      5.6431     52.0    5.817352  ...    2.547945     37.85     -122.25
4      3.8462     52.0    6.281853  ...    2.181467     37.85     -122.25
...      ...      ...      ...      ...      ...      ...
20635  1.5603     25.0    5.045455  ...    2.560606     39.48     -121.09
20636  2.5568     18.0    6.114035  ...    3.122807     39.49     -121.21
20637  1.7000     17.0    5.205543  ...    2.325635     39.43     -121.22
20638  1.8672     18.0    5.329513  ...    2.123209     39.43     -121.32
20639  2.3886     16.0    5.254717  ...    2.616981     39.37     -121.24

[20640 rows x 8 columns]
Mean Squared Error: 1.2923314440807299
R-squared: 0.013795337532284901
Intercept: 1.654762268596842
Coefficient: [0.07675559]
Mean Squared Error: 0.5558915986952441
R-squared: 0.575787706032451
Intercept: -37.02327770606414
Coefficient: [ 4.48674910e-01  9.72425752e-03 -1.23323343e-01  7.83144907e-01
 -2.02962058e-06 -3.52631849e-03 -4.19792487e-01 -4.33708065e-01]

```

PRACTICAL NO. – 7**AIM:- Logistic Regression and Decision Tree SOURCE CODE:-**

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import warnings
from scipy import stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd
from sklearn.datasets import fetch_california_housing, load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score, precision_score, recall_score, classification_report

# Suppress warnings
warnings.filterwarnings('ignore')

# Load dataset
df = sb.load_dataset('mpg')
print(df)
```



```
# Describe horsepower and model year columns
print(df['horsepower'].describe()) print(df['model_year'].describe())

# Categorize horsepower into bins bins =
[0, 75, 150, 240]
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm', 'h'])
c = df['horsepower_new'] print(c)

# Categorize model year into bins ybins =
[69, 72, 74, 84]
labels = ['t1', 't2', 't3']
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins,
labels=labels)
newyear = df['modelyear_new']
print(newyear)

# Create a contingency table
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new']) print(df_chi)
```

```
# Perform chi-square test
chi2_stat, p_value, dof, expected = stats.chi2_contingency(df_chi)
print(f'Chi-Square Statistic: {chi2_stat}')
print(f'P-value: {p_value}') print(f'Degrees of
Freedom: {dof}')
print(f'Expected Frequencies:\n{expected}')

# Define groups for ANOVA
group1 = [23, 25, 29, 34, 30]
group2 = [19, 20, 22, 24, 25]
group3 = [15, 18, 20, 21, 17]
group4 = [28, 24, 26, 30, 29]

# Combine data into a DataFrame
data = pd.DataFrame({'value': group1 + group2 + group3 + group4, 'group':
                    ['Group1'] * len(group1) + ['Group2'] *
len(group2) +
                    ['Group3'] * len(group3) + ['Group4'] *
len(group4)})

# Perform one-way ANOVA
f_statistics, p_value_anova = stats.f_oneway(group1, group2, group3, group4)
print("\nOne-way ANOVA:")
print("F-statistics:", f_statistics)
```

```
print("P-value:", p_value_anova)
```

```
# Perform Tukey-Kramer post-hoc test
```

```
tukey_results = pairwise_tukeyhsd(data['value'], data['group'])
```

```
print("\nTukey-Kramer post-hoc test:")
```

```
print(tukey_results)
```

```
# Load California housing dataset housing =
```

```
fetch_california_housing()
```

```
housing_df = pd.DataFrame(housing.data,  
columns=housing.feature_names)
```

```
print(housing_df)
```

```
# Add target variable housing_df['PRICE'] =  
housing.target
```

```
# Multiple Linear Regression
```

```
X = housing_df.drop('PRICE', axis=1) y =  
housing_df['PRICE']
```

```
# Split dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train linear regression model model =  
LinearRegression() model.fit(X_train,  
y_train)  
  
# Make predictions  
y_pred = model.predict(X_test)  
  
# Evaluate model  
mse = mean_squared_error(y_test, y_pred) r2 =  
r2_score(y_test, y_pred)  
print("Mean Squared Error:", mse)  
print("R-squared:", r2) print("Intercept:",  
model.intercept_) print("Coefficient:",  
model.coef_)  
  
# Load the Iris dataset and classification problem iris =  
load_iris()  
iris_df = pd.DataFrame(data=np.c_[iris['data'], iris['target']],  
columns=iris['feature_names'] + ['target'])  
binary_df = iris_df[iris_df['target'] != 2] X =  
binary_df.drop('target', axis=1)  
y = binary_df['target']  
  
# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train a logistic regression model and evaluate its performance
```

```
logistic_model = LogisticRegression() logistic_model.fit(X_train,  
y_train)
```

```
y_pred_logistic = logistic_model.predict(X_test)
```

```
print("\nLogistic Regression Metrics")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred_logistic))
```

```
print("Precision:", precision_score(y_test, y_pred_logistic))
```

```
print("Recall:", recall_score(y_test, y_pred_logistic))
```

```
print("\nClassification Report") print(classification_report(y_test,  
y_pred_logistic))
```

```
# Train a decision tree model and evaluate its performance decision_tree_model =
```

```
DecisionTreeClassifier() decision_tree_model.fit(X_train, y_train)
```

```
y_pred_tree = decision_tree_model.predict(X_test)
```

```
print("\nDecision Tree Metrics")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred_tree))
```

```
print("Precision:", precision_score(y_test, y_pred_tree))
```

```
print("Recall:", recall_score(y_test, y_pred_tree))
```

```
print("\nClassification Report")  
print(classification_report(y_test, y_pred_tree))
```

OUTPUT:-

```
Logistic Regression Metrics  
Accuracy: 1.0  
Precision: 1.0  
Recall: 1.0  
  
Classification Report  
              precision    recall  f1-score   support  
  
    0.0         1.00      1.00      1.00        12  
    1.0         1.00      1.00      1.00         8  
  
   accuracy                1.00          20  
  macro avg              1.00      1.00      1.00          20  
weighted avg              1.00      1.00      1.00          20  
  
Decision Tree Metrics  
Accuracy: 1.0  
Precision: 1.0  
Recall: 1.0
```

PRACTICAL NO. – 8**AIM:- K-MEANS CLUSTERING SOURCE****CODE:-**

```
import pandas as pd

from sklearn.preprocessing import MinMaxScaler
from sklearn.cluster import KMeans

import matplotlib.pyplot as plt


# Load dataset data

=
pd.read_csv("C:\\Users\\Reape\\Downloads\\wholesale\\wholesale.csv ")
data.head()


# Define categorical and continuous features categorical_features =
['Channel', 'Region']

continuous_features = ['Fresh', 'Milk', 'Grocery', 'Frozen',
'Detergents_Paper', 'Delicassen']

data[continuous_features].describe()


# Convert categorical variables into dummy variables for col in
categorical_features:

    dummies = pd.get_dummies(data[col], prefix=col) data =
    pd.concat([data, dummies], axis=1)
```

```
data.drop(col, axis=1, inplace=True)
```

```
data.head()
```

```
# Scale the data
```

```
mms = MinMaxScaler()
```

```
mms.fit(data)
```

```
data_transformed = mms.transform(data)
```

```
# Elbow Method to find optimal k
```

```
sum_of_squared_distances = []
```

```
K = range(1, 15)
```

```
for k in K:
```

```
    km = KMeans(n_clusters=k) km =
```

```
    km.fit(data_transformed)
```

```
    sum_of_squared_distances.append(km.inertia_)
```

```
# Plot the Elbow Method
```


```
plt.plot(K, sum_of_squared_distances, 'bx-')
```

```
plt.xlabel('k')
```

```
plt.ylabel('Sum of Squared Distances')
```

```
plt.title('Elbow Method for Optimal k')
```

```
plt.show()
```


OUTPUT:-

PRACTICAL NO. 9**AIM:- Principal Component Analysis (PCA) SOURCE CODE:-**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

# Load the Iris dataset iris
iris = load_iris()

iris_df = pd.DataFrame(data=np.c_[iris['data'], iris['target']],
                       columns=iris['feature_names'] + ['target'])

# Separate features and target
X = iris_df.drop('target', axis=1)
y = iris_df['target']

# Standardize the features scaler
scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)

# Apply PCA
```

```
pca = PCA()
X_pca = pca.fit_transform(X_scaled) explained_variance_ratio =
pca.explained_variance_ratio_

# Plot cumulative explained variance
plt.figure(figsize=(8, 6))
plt.plot(np.cumsum(explained_variance_ratio), marker='o',
linestyle='--')
plt.title('Explained Variance Ratio') plt.xlabel('Number of
Principal Components') plt.ylabel('Cumulative Explained
Variance Ratio') plt.grid(True)
plt.show()

# Determine the number of components to explain 95% variance
cumulative_variance_ratio = np.cumsum(explained_variance_ratio) n_components =
np.argmax(cumulative_variance_ratio >= 0.95) + 1
print(f"Number of principal components to explain 95% variance:
{n_components}")

# Reduce dimensions using selected number of components pca =
PCA(n_components=n_components)
X_reduced = pca.fit_transform(X_scaled)
```

```
# Scatter plot of the reduced data
```

```
plt.figure(figsize=(8, 6))
```


```
plt.scatter(X_reduced[:, 0], X_reduced[:, 1], c=y, cmap='viridis', s=50, alpha=0.5)
```


```
plt.title('Data in Reduced-dimensional Space') plt.xlabel('Principal  
Component 1')
```

```
plt.ylabel('Principal Component 2')
```

```
plt.colorbar(label='Target') plt.show()
```

OUTPUT:-





----- RESTART: E./all notes/DS/PL
Number of principal components to explain 95% variance: 2

PRACTICAL NO. – 10**AIM:-Data Visualization and Storytelling SOURCE CODE:-**

```
import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

import numpy as np

# Generate random data

np.random.seed(42) # Set a seed for reproducibility

# Create a DataFrame with random data data
= pd.DataFrame({
    'variable1': np.random.normal(0, 1, 1000),
    'variable2': np.random.normal(2, 2, 1000) + 0.5 *
np.random.normal(0, 1, 1000),
    'variable3': np.random.normal(-1, 1.5, 1000),
    'category': pd.Series(np.random.choice(['A', 'B', 'C', 'D'], size=1000, p=[0.4,
0.3, 0.2, 0.1]),
                        dtype='category')
})

# Create a scatter plot to visualize the relationship between two variables
```

```
plt.figure(figsize=(10, 6))  
plt.scatter(data['variable1'], data['variable2'], alpha=0.5) plt.title('Relationship  
between Variable 1 and Variable 2', fontsize=16) plt.xlabel('Variable 1',  
fontsize=14)  
plt.ylabel('Variable 2', fontsize=14) plt.show()
```

Create a bar chart to visualize the distribution of a categorical variable

```
plt.figure(figsize=(10, 6)) sns.countplot(x='category',  
data=data) plt.title('Distribution of Categories',  
fontsize=16) plt.xlabel('Category', fontsize=14)  
plt.ylabel('Count', fontsize=14) plt.xticks(rotation=45)  
plt.show()
```

Create a heatmap to visualize the correlation between numerical variables

```
plt.figure(figsize=(10, 8))  
numerical_cols = ['variable1', 'variable2', 'variable3']  
sns.heatmap(data[numerical_cols].corr(), annot=True,  
cmap='coolwarm')  
plt.title('Correlation Heatmap', fontsize=16)  
plt.show()
```

```
# Data Storytelling
```

```
print("Title: Exploring the Relationship between Variable 1 and Variable 2")
```

```
print("\nThe scatter plot (Figure 1) shows the relationship between Variable 1  
and Variable 2.")
```

```
print("\nScatter Plot")
```

```
print("Figure 1: Scatter Plot of Variable 1 and Variable 2")
```

```
print("\nTo better understand the distribution of the categorical variable  
'category', we created a ")
```

```
print("\nBar Chart")
```



```
print("Figure 2: Distribution of Categories")
```


```
print("\nAdditionally, we explored the correlation between numerical variables using a  
heatmap")
```

```
print("\nHeatmap")
```

```
print("Figure 3: Correlation Heatmap")
```

```
print("\nIn summary, the visualizations and analysis provide insights into the  
relationships")
```


OUTPUT:-



= RESTART: D:/pract 10.py

Title: Exploring the Relationship between Variable 1 and Variable 2

The scatter plot (Figure 1) shows the relationship between Variable 1 and Variable 2.

Scatter Plot

Figure 1: Scatter Plot of Variable 1 and Variable 2

To better understand the distribution of the categorical variable 'category', we created a

Bar Chart

Figure 2: Distribution of Categories

Additionally, we explored the correlation between numerical variables using a heatmap

Heatmap

Figure 3: Correlation Heatmap

In summary, the visualizations and analysis provide insights into the relationships